

Overlap Savings Analysis for Real-Time Voice Translation

Document Purpose: Explain why pipeline overlap optimization provides zero benefit for real-time streaming, and identify potential architectural improvements.

Target Audience: AI systems analyzing optimization opportunities

Last Updated: 2026-02-19 21:32 HKT (Rev 3 - Week 0 Added)

1. Executive Summary

The current parallel pipeline (`orchestrator_parallel.py`) implements worker threads for VAD, ASR, Translation, and Output. However, **overlap savings = 0ms** in real-time streaming mode because the pipeline is fundamentally **I/O bound by human speech speed**, not **compute bound**.

Key Insight: Human speech (3-8 seconds per segment) is 5-10x slower than processing time (~0.7 seconds). When ASR finishes segment N, segment N+1 hasn't been fully spoken yet.

2. Timing Analysis

2.1 Real-Time Streaming Timeline

```
gantt
    title Real-Time Streaming: Sequential by Nature
    dateFormat X
    axisFormat %s

    section Human Speech
    Segment 1 (5s speech)      :active, speech1, 0, 5
    Silence (400ms)            :crit, silence1, 5, 5.4
    Segment 2 (4s speech)      :active, speech2, 5.4, 9.4

    section VAD Detection
    VAD: Seg 1 detected        :done, vad1, 5, 5.1
    VAD: Seg 2 detected        :done, vad2, 9.4, 9.5

    section ASR Processing
    ASR: Process Seg 1 (~450ms) :active, asr1, 5.1, 5.55
    ASR: WAITING (idle)         :crit, asr_wait, 5.55, 9.5
    ASR: Process Seg 2 (~450ms) :active, asr2, 9.5, 9.95

    section Translation
    MT: Translate Seg 1 (~250ms):active, mt1, 5.55, 5.8
    MT: WAITING (idle)          :crit, mt_wait, 5.8, 9.95
    MT: Translate Seg 2 (~250ms):active, mt2, 9.95, 10.2
```

Observation: Notice the large **idle gaps** (red segments). When ASR finishes at T=5.55s, the next segment hasn't even started being spoken (starts at T=5.4s, but VAD needs silence to detect it at T=9.4s).

2.2 Batch Processing Timeline (For Comparison)

```

gantt
    title Batch Processing: True Parallelism Possible
    dateFormat X
    axisFormat %s

    section VAD (File Scan)
    Detect ALL segments      :done, vad_batch, 0, 2

    section ASR Worker Pool
    ASR: Segment 1          :active, asr_b1, 2, 2.45
    ASR: Segment 2          :active, asr_b2, 2.45, 2.9
    ASR: Segment 3          :active, asr_b3, 2.9, 3.35

    section Translation Worker Pool
    MT: Segment 1           :active, mt_b1, 2.45, 2.7
    MT: Segment 2           :active, mt_b2, 2.7, 2.95
    MT: Segment 3           :active, mt_b3, 2.95, 3.2

    section Output
    Output Seg 1            :done, out_b1, 2.7, 2.8
    Output Seg 2            :done, out_b2, 2.95, 3.05
    Output Seg 3            :done, out_b3, 3.2, 3.3

```

Observation: In batch mode, all segments are available at T=0 (after VAD scan), enabling true pipeline parallelism.

3. Mathematical Constraint

3.1 Timing Parameters (from STATUS.md)

Parameter	Symbol	Value	Description
Speech Duration	T_speech	3-8s	Human speech segment length
Silence Duration	T_silence	400ms	VAD silence threshold
ASR Processing	T_asr	~450ms	faster-whisper base model
Translation	T_mt	~250ms	MarianMT inference
Total Processing	T_process	~700ms	ASR + MT

3.2 Overlap Condition

For pipeline overlap to provide savings:

Next segment must be available before current processing finishes

$T_{available} < T_{process}$

Real-time case:

$T_{available} = T_{speech} + T_{silence} = 3\text{--}8\text{s} + 0.4\text{s} = 3.4\text{--}8.4\text{s}$
 $T_{process} = 0.7\text{s}$

$3.4\text{--}8.4\text{s} < 0.7\text{s} = \text{FALSE } \times$

Batch case:

$T_{available} = 0$ (all segments ready immediately)
 $T_{process} = 0.7\text{s}$

$0 < 0.7\text{s} = \text{TRUE } \checkmark$

3.3 Worker Utilization

```
pie title Real-Time Worker Utilization
    "ASR Worker Idle (Waiting for Audio)" : 85
    "ASR Worker Active (Processing)" : 15
```

```
pie title Batch Worker Utilization
    "ASR Worker Idle" : 5
    "ASR Worker Active" : 95
```

4. Current Architecture Deep Dive

4.1 Parallel Pipeline Structure

```
flowchart TB
    subgraph AudioCapture ["Audio Capture Thread"]
        A[Microphone Input] --> B[Ring Buffer]
    end

    subgraph VADWorker ["VAD Worker (1 thread)"]
        B --> C{Speech Detected?}
        C -->|Yes| D[Buffer Speech Segment]
```

```

C -->|No| C
D --> E[Queue to ASR]
end

subgraph ASRWorker["ASR Worker (1 thread)"]
    E --> F[ASR Inference]
    F -->|2-thread pool| G[CTranslate2]
    G --> H[Queue to MT]
end

subgraph MTWorker["Translation Worker (1 thread)"]
    H --> I[MT Inference]
    I -->|2-thread pool| J[MarianMT/NLLB]
    J --> K[Queue to Output]
end

subgraph OutputWorker["Output Worker (1 thread)"]
    K --> L[Display/Log]
end

style VADWorker fill:#ffcccc
style ASRWorker fill:#ffcccc
style MTWorker fill:#ffcccc
style OutputWorker fill:#ffcccc

```

4.2 The Blocking Problem

```

sequenceDiagram
    participant H as Human
    participant V as VAD Worker
    participant A as ASR Worker
    participant T as Translation Worker
    participant O as Output Worker

    Note over H,O: Segment 1 Begins

    loop Speech Duration (5s)
        H->>V: Audio chunks
        V->>V: Accumulating...
    end

    H->>V: Silence (400ms)
    V->>V: Detect end of speech
    V->>A: Segment 1 ready

    rect rgb(200, 255, 200)
        Note over A: ASR Active (450ms)
        A->>A: Process audio
    end
    A->>T: ASR result

```

```

rect rgb(200, 255, 200)
    Note over T: MT Active (250ms)
    T->>T: Translate text
end
T->>0: Translation result
0->>0: Output

Note over H,0: ASR/MT Workers IDLE waiting for Segment 2

rect rgb(255, 200, 200)
    Note over A,T: IDLE (~3-7s)
Next segment not yet spoken
end

loop Speech Duration (4s)
    H->>V: Audio chunks
    V->>V: Accumulating...
end

H->>V: Silence
V->>A: Segment 2 ready
Note over A: ASR starts again...

```

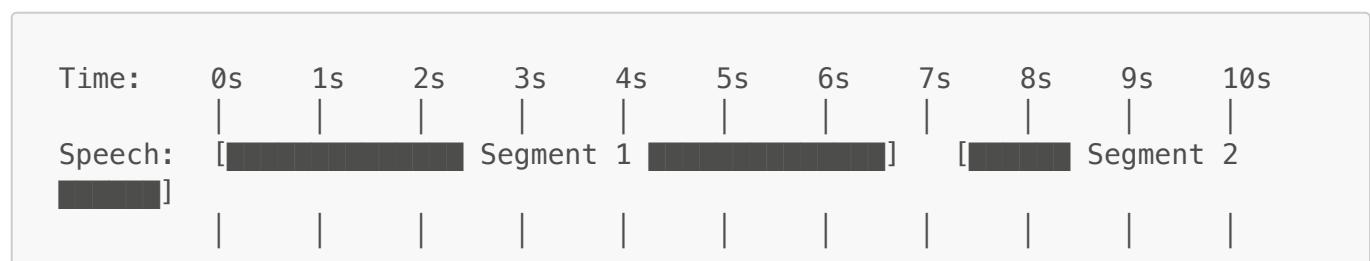
Critical Observation: The 4 workers run sequentially in practice because each waits for the previous stage to produce data, and data production is bound by human speech speed.

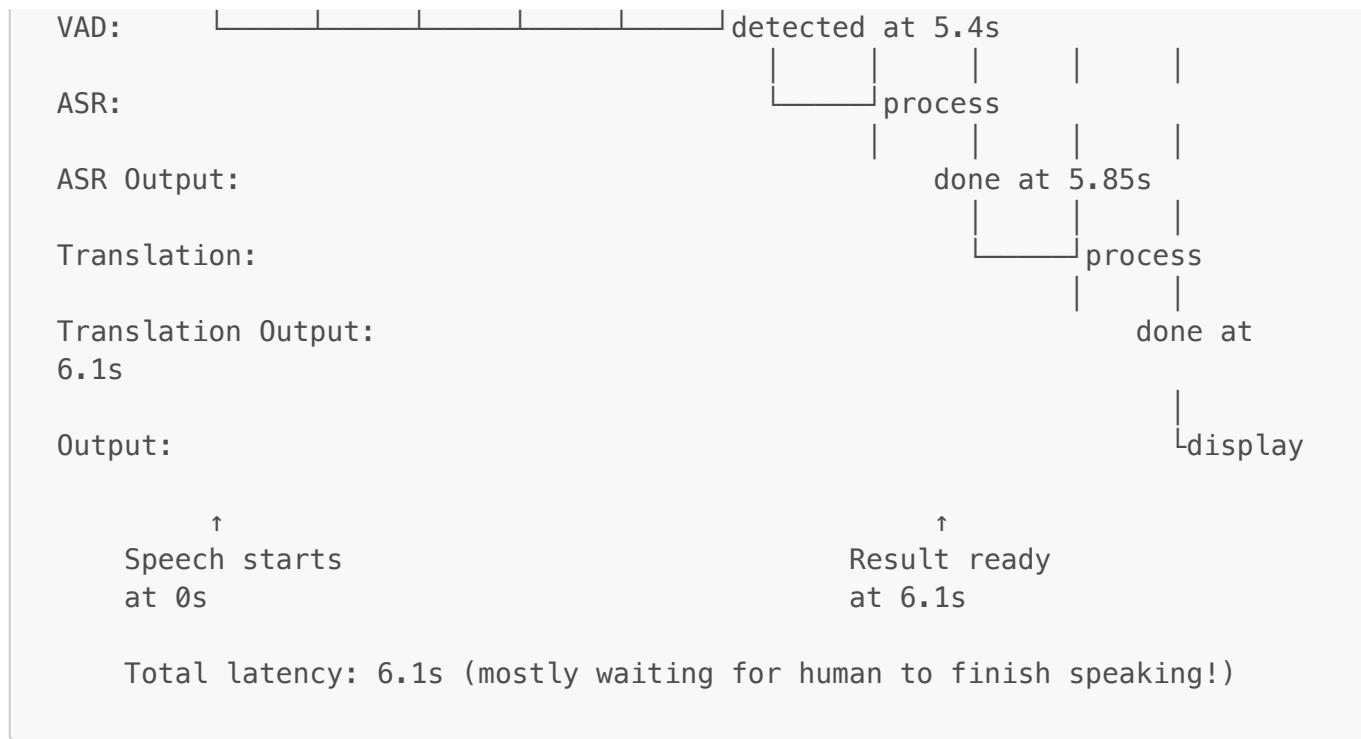
5. Fundamental Limitations

5.1 Why Current Parallel Architecture Can't Help

Limitation	Explanation
Sequential Data Dependency	VAD must wait for speech+silence before producing a segment. ASR must wait for VAD. MT must wait for ASR.
No Lookahead	In real-time, you cannot see future audio. Batch can scan the entire file first.
Speech-Processing Speed Mismatch	Humans speak at ~150 words/min. ASR processes at ~1000 words/min (after receipt).
VAD as Bottleneck	VAD cannot detect segments faster than they are spoken.

5.2 The "Idle Gap" Problem





6. Potential Improvement Directions

6.1 Streaming/Incremental ASR

Concept: Process audio incrementally as it arrives, not waiting for segment end.

```

gantt
    title Streaming ASR: Potential for Overlap
    dateFormat X
    axisFormat %s

    section Human Speech
    Word 1 "Hello"           :active, w1, 0, 0.5
    Word 2 "world"           :active, w2, 0.6, 1.1
    Word 3 "today"           :active, w3, 1.2, 1.7

    section Streaming ASR
    Process Word 1           :done, sw1, 0.5, 0.7
    Process Word 2           :done, sw2, 1.1, 1.3
    Process Word 3           :done, sw3, 1.7, 1.9

    section Speculative Translation
    Start MT (partial)      :crit, mt_p1, 0.7, 0.9
    Update MT (partial)     :crit, mt_p2, 1.3, 1.5
    Finalize MT              :active, mt_f, 1.9, 2.1

```

Challenges:

- ASR accuracy on partial audio (context loss)
- MT quality on incomplete sentences

- Stitching partial results

Implementation: faster-whisper supports `word_timestamps=True` and incremental decoding.

6.2 Sentence-Breaking VAD (Chunk-Based)

Concept: Break long sentences into smaller chunks (e.g., 2-second chunks) instead of waiting for full silence.

```
flowchart LR
    A[2s chunk 1] -->|ASR| B[Partial result]
    C[2s chunk 2] -->|ASR| D[Partial result]
    E[2s chunk 3] -->|ASR| F[Partial result]
    B --> G{Stitch}
    D --> G
    F --> G
    G --> H[Full sentence]
    H --> I[MT]
```

Challenges:

- Context window management for ASR
- Sentence boundary detection without full pause
- Overlapping chunk processing

6.3 Predictive Pre-processing

Concept: Use lightweight model to predict if speech is ending soon, start ASR early.

```
sequenceDiagram
    participant V as VAD + Predictor
    participant A as ASR
    participant T as Translation

    Note over V: Speech detected for 3s
    V->>V: Lightweight predictor:
    "90% probability speech ends in 500ms"
    V->>A: Start ASR early (speculative)

    alt Speech actually ends
        A->>A: Result ready immediately!
        A->>T: Translation starts
    else Speech continues
        A->>A: Discard partial result
        (wasted compute)
        Note over A: Trade-off:
    Latency vs Compute
    end
```

Challenges:

- Prediction accuracy
- Wasted compute on false predictions
- Model size (must be very lightweight)

6.4 Pipeline Stage Interleaving

Concept: Process multiple speakers or audio sources (if available) to keep workers busy.

```

gantt
    title Multi-Source Interleaving
    dateFormat X
    axisFormat %s

    section Source A
    Speech A1           :active, sa1, 0, 5

    section Source B
    Speech B1           :active, sb1, 2, 7

    section ASR Worker
    Process A1          :active, asra1, 5, 5.45
    Process B1          :active, asrb1, 7, 7.45

    section Translation Worker
    Translate A1        :active, mta1, 5.45, 5.7
    Translate B1        :active, mtb1, 7.45, 7.7

```

Note: This only works for multi-source scenarios (e.g., conference calls), not single-speaker translation.

7. Key Questions for Analysis

When considering improvements, answer these:

1. Is latency actually a problem?

- Current end-to-end: ~700-850ms after speech ends
- Human perception threshold: ~200-300ms feels "instant"
- But most latency is speech duration itself, not processing

2. What's the real bottleneck?

- Compute: No (processing is 5-10x faster than speech)
- I/O: Yes (waiting for audio)
- Memory: No (models already loaded)

3. Can we reduce speech duration dependency?

- Only by processing partial segments

- Trade-off: Quality vs Latency

4. Is batch mode the actual use case?

- If translating files/videos, overlap optimization is valuable
 - If translating live speech, current architecture is near-optimal
-

8. Conclusion

The overlap savings = 0ms is architecturally correct for real-time streaming.

The parallel pipeline ([orchestrator_parallel.py](#)) is designed for throughput optimization (batch processing), but real-time streaming is fundamentally limited by the speed of human speech production.

To achieve actual latency improvements in real-time mode, you would need to:

1. Process audio **before** the speaker finishes (streaming ASR)
2. Accept **quality trade-offs** (partial context)
3. Implement **speculative execution** (risk wasted compute)

For the current single-speaker real-time use case, the architecture is already near-optimal. The parallel pipeline provides value for batch video/file processing, where overlap savings are real and significant.

9. References

- Current Pipeline: [src/core/pipeline/orchestrator_parallel.py](#)
 - Performance Metrics: [STATUS.md](#) (Section "Performance Metrics")
 - ASR Implementation: [src/core/asr/faster_whisper.py](#)
 - VAD Implementation: [src/audio/vad/silero_vad_adaptive.py](#)
-

10. Update: Streaming Solution Design (REVISED)

After evaluation of architectural suggestions (see [docs/evaluation_streaming_suggestions.md](#)), a **Hybrid Streaming Mode with Partial Translation** has been designed.

10.1 The Solution: Hybrid Draft/Final Mode

```
Current (Sequential – No Meaning Until End):
Speech: [ ] (5s) → Silence → ASR → MT → "Hola mundo"
at 5.8s

Hybrid (Overlapped – Meaning Early):
Speech: [ ] (5s)
Draft ASR 1: [ASR] → "Hello world"
Draft MT 1: [MT] → "Hola mundo" at 2.75s ✓ MEANING!
Draft ASR 2: [ASR] → "Hello world today"
Draft MT 2: [MT] → "Hola mundo hoy" at 4.5s
```

Final:
at 5.3s

[ASR+MT] → "Hola mundo hoy"

10.2 REVISED Key Design Decisions

Aspect	Original	REVISED	Rationale
Draft interval	Every 2s	Adaptive (every 2s, skip if paused)	Reduce compute overhead
Draft confidence	Low (0.6)	Low (0.6) ✓	Show grey/italic
Draft translation	No (ASR only)	Yes (conditional on verb/punctuation)	Users need meaning , not just words
Compute strategy	Standard	INT8 for drafts, standard for final	Manage 3x ASR overhead
Context window	Chunk-based	Cumulative (0-N, not just N-2→N)	Ensures grammatical consistency
UI transition	Replace	Diff-based with highlight	Smooth, show what changed

10.3 Critical Improvement: Draft Translation with Semantic Gating

Why Conditional Translation Matters:

Approach	T=2s	T=5.3s	User Experience
Without Translation	"Hello world..." (EN)	"Hola mundo" (ES)	Wait 5.3s for meaning ✗
With Translation	"Hola mundo..." (ES)	"Hola mundo" (ES)	Get meaning at 2s ✓

Semantic Gating Rules (only translate if complete thought):

```
def should_translate_draft(text, source_lang, target_lang):
    # Must have minimum content
    if len(text.split()) < 2:
        return False

    # Must have verb or punctuation (complete thought)
    has_verb = any(v in text.lower() for v in VERBS[source_lang])
    has_punct = any(text.endswith(p) for p in ['.', '!', '?', '。'])

    # Word order safety for SOV languages
    # SOV languages (Japanese, Korean) put verb at END
    # Translating before verb arrives = grammatical chaos
    SOV_LANGUAGES = ['ja', 'ko', 'de', 'tr']
    if target_lang in SOV_LANGUAGES:
        if not has_punct:
            return False # Must wait for sentence end
```

```

        return has_verb or has_punct

# Examples:
# "Hello" → No translation (incomplete)
# "Hello world" → No translation (no verb)
# "I went" → Translate! (has verb)
# "Hello world." → Translate! (has punctuation)
# EN→JA: "I went to" → No (wait for Japanese verb at end)
# EN→JA: "I went to store." → Translate! (has punctuation)

```

Language-Specific Rules:

Language Pair	Draft Translation	Notes
EN→ES, EN→FR	✓ Yes	SVO→SVO (similar order)
EN→JA, EN→KO	⚠ Conditional	SVO→SOV (wait for punctuation)
JA→EN	✓ Yes	SOV→SVO (can translate early)

10.4 Managing 3x Compute Overhead

Problem: 5s segment → Draft at 2s + Draft at 4s + Final at 5s = **3x ASR calls**

Mitigation Strategy:

Technique	Implementation	Impact	Priority
INT8 Quantization	Draft: <code>compute_type="int8"</code>	2x faster drafts	Week 2
Adaptive Skipping	Skip draft if VAD detects pause >500ms	Reduce unnecessary calls	Week 2
Queue Monitoring	Skip draft if queue depth > 2; Alert on overflow	Prevent backlog, fix data loss	🔴 Week 0
Cumulative Context	Draft 2 reuses Draft 1 computation conceptually	Better quality per compute	Week 2
Error Logging	No silent failures, comprehensive tracing	Debug sentence loss	🔴 Week 0
Platform Testing	Intel i7 (OpenVINO), Mac M1 (CoreML)	Optimize INT8 per platform	🔴 Week 0

Compute Estimate:

```

Unoptimized: 3 × 450ms = 1350ms (3x overhead)
With INT8:   2 × 225ms + 450ms = 900ms (2x overhead)
With Adaptive: ~1.5x average (depends on speech pattern)

```

10.5 Expected Improvements

Metric	Current	Target	Priority	Notes
Sentence Loss Rate	Bug exists (~?)	0%	● Week 0 Critical	Data integrity first
TTFT (any output)	~5000ms	< 2000ms	● Week 2	Draft visible
Meaning Latency (translated)	~5000ms	< 2000ms	● Week 2	★ Critical improvement
Ear-to-Voice Lag	~700ms	< 500ms	● Week 3	Final optimization
Compute Overhead	1x	~1.5x	● Week 2	Managed with INT8 + adaptive

10.6 Implementation Plan (REVISED with Week 0)

See detailed plan: [docs/design/streaming_latency_optimization_plan.md](#)

Week 0: CRITICAL - Fix Sentence Loss Bug ●

MUST complete before any optimization!

Task	Deliverable
Add segment sequence tracking	UUID per segment, trace capture → VAD → ASR → MT → Output
Add queue depth monitoring	Alert if queue > 3 segments
Add comprehensive error logging	No silent failures anywhere
Stress test: 10-min continuous speech	Verify 0% sentence loss
Fix root cause	Queue overflow? VAD threshold? Race condition?
Platform-specific testing	Intel i7: OpenVINO INT8, Mac M1: CoreML optimization

Week 1: Metrics + Adaptive Config (TTFT/Lag, segment=4s, adaptive controller)

Week 2: StreamingASR (cumulative context, INT8 drafts) + StreamingTranslator (semantic gating)

Week 3: Diff-based UI (smooth transitions, stability indicators) + Integration

10.7 What Was Rejected

Approach	Reason
Wait-k Translation	MarianMT doesn't support streaming; fails for different word order languages
AsyncIO Rewrite	Over-engineering; ThreadPool performs equally for ML tasks

Approach	Reason
True Streaming ASR	Requires model change (RNN-T); too much effort
Draft Without Translation	Provides "listening feedback" but not "meaning" - insufficient UX

Document written for AI analysis and improvement planning.

Updated 2026-02-19 with REVISED streaming solution (partial translation, semantic gating, INT8 optimization).